

# Deng\_Yehong\_HW7

March 15, 2020

```
[29]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

1. (5 points) Imitate the k-means random initialization part of the algorithm by assigning each observation to a cluster at random.

```
[30]: input_1 = np.array([5,8,7,8,3,4,2,3,4,5])
input_2 = np.array([8,6,5,4,3,2,2,8,9,8])
```

```
[31]: np.random.seed(456123)
df = pd.DataFrame({'input_1': input_1, 'input_2': input_2})
get_labels = np.random.choice(3, 10, replace = True)
df['k'] = get_labels
df
```

```
[31]:
```

	input_1	input_2	k
0	5	8	1
1	8	6	0
2	7	5	2
3	8	4	2
4	3	3	0
5	4	2	1
6	2	2	0
7	3	8	0
8	4	9	2
9	5	8	1

2. (5 points) Compute the cluster centroid and update cluster assignments for each observation iteratively based on spatial similarity.

```
[32]: df_k3 = df.copy()
      for i in range(100):
          centroids = {}
          for j in range(3):
              cen_input1 = df_k3[df_k3['k'] == j]['input_1'].mean()
              cen_input2 = df_k3[df_k3['k'] == j]['input_2'].mean()
              centroids[j] = (cen_input1, cen_input2)
          updated_k = []
          for index, row in df_k3.iterrows():
              min_distance = 100
              for key, value in centroids.items():
                  distance = ((row['input_1'] - value[0]) ** 2 + (row['input_2'] -
→value[1])**2)**0.5
                  if distance < min_distance:
                      min_distance = distance
                      update_k = key
              updated_k.append(update_k)
          df_k3['k'] = updated_k
df_k3
```


```
[32]:
```

	input_1	input_2	k
0	5	8	1
1	8	6	2
2	7	5	2
3	8	4	2
4	3	3	0
5	4	2	0
6	2	2	0
7	3	8	1
8	4	9	1
9	5	8	1

3. (5 points) Present a visual description of the final, converged (stopped) cluster assignments.

```
[33]: sns.scatterplot(x = 'input_1', y = 'input_2', hue = 'k', palette = ['blue',
→'green', 'red'], data = df_k3)
```

```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1f2ba6ca648>
```



4. (5 points) Now, repeat the process, but this time initialize at  $k = 2$  and present a final cluster assignment visually next to the previous search at  $k = 3$ .

```
[35]: df2 = pd.DataFrame({'input_1': input_1, 'input_2': input_2})
get_labels2 = np.random.choice(2, 10, replace = True)
df2['k'] = get_labels2
df2
```

```
[35]:
```

	input_1	input_2	k
0	5	8	0
1	8	6	0
2	7	5	1
3	8	4	0
4	3	3	1
5	4	2	1
6	2	2	1
7	3	8	0
8	4	9	1
9	5	8	0

```
[36]: df_k2 = df2.copy()
for i in range(100):
    centroids = {}
    for j in range(2):
        cen_input1 = df_k2[df_k2['k'] == j]['input_1'].mean()
```

```

cen_input2 = df_k2[df_k2['k'] == j]['input_2'].mean()
centroids[j] = (cen_input1, cen_input2)
updated_k = []
for index, row in df_k2.iterrows():
    min_distance = 100
    for key, value in centroids.items():
        distance = ((row['input_1'] - value[0]) ** 2 + (row['input_2'] -
→value[1])**2)**0.5
        if distance < min_distance:
            min_distance = distance
            update_k = key
    updated_k.append(update_k)
df_k2['k'] = updated_k
df_k2

```

```

[36]:
input_1  input_2  k
0         5         8  0
1         8         6  0
2         7         5  0
3         8         4  0
4         3         3  1
5         4         2  1
6         2         2  1
7         3         8  0
8         4         9  0
9         5         8  0

```

```

[38]: #subplots referenced from https://matplotlib.org/api/_as_gen/matplotlib.pyplot.
      →subplots.html
fig, axs = plt.subplots(1,2)
color = np.array(['blue', 'green', 'red'])
axs[0].scatter(df_k3['input_1'], df_k3['input_2'], c = color[df_k3['k']])
axs[0].set_xlabel('input_1')
axs[0].set_ylabel('input_2')
axs[0].set_title('k = 3 Clustering')


axs[1].scatter(df_k2['input_1'], df_k2['input_2'], c = color[df_k2['k']])
axs[1].set_xlabel('input_1')
axs[1].set_ylabel('input_2')
axs[1].set_title('k = 2 Clustering')

```

```

[38]: Text(0.5, 1.0, 'k = 2 Clustering')

```



5. (10 points) Did your initial hunch of 3 clusters pan out, or would other values of  $k$ , like 2, fit these data better? Why or why not?

According to the  $k=3$  clustering plot, the red cluster located in the middle between the green and blue ones. The distance between red and green and that between red and blue should be very similar. Hence, it make sense that the graph should have three clusters. Nonetheless, in the the  $k = 2$  clustering graph, the original green and red clusters are now clustered together. As a result, the distance among points within each cluster in the right graph is higher. Therefore,  $k = 3$  performs better than  $k = 2$ .

6. (15 points) Perform PCA on the dataset and plot the observations on the first and second principal components. Describe your results, e.g.,

What variables appear strongly correlated on the first principal component?

What about the second principal component?

```
[49]: df_wiki = pd.read_csv('wiki.csv')
predictors = df_wiki.columns
df_wiki = StandardScaler().fit_transform(df_wiki)
pca = PCA(n_components = 2, random_state = 456123)
df_pca = pd.DataFrame(data = pca.fit(df_wiki).components_.T, index = _
    ↪predictors, columns = ['PC1', 'PC2'])
```

```
[50]: df_pca.nlargest(5, 'PC1')
```

```
[50]:
```

	PC1	PC2
bi2	0.230924	0.083430
bi1	0.226193	0.056372
use3	0.218809	0.155158
use4	0.214558	0.160868
pu3	0.210863	0.028776

```
[51]: df_pca.nsmallest(5, 'PC2')
```

```
[51]:
```

	PC1	PC2
peu1	0.061228	-0.271740
inc1	0.104667	-0.245440
sa3	0.120376	-0.242315
sa1	0.121658	-0.229925
enj2	0.131110	-0.227596

```
[52]: pca = PCA(n_components = 2, random_state = 456123)
PC = pca.fit_transform(df_wiki)
df_PC = pd.DataFrame(data = PC, columns = ['PC1', 'PC2'])
df_PC
```


```
[52]:
```

	PC1	PC2
0	-0.150216	-1.983507
1	-3.314020	-0.791893
2	-4.682484	-0.312228
3	1.774200	1.986370
4	7.254695	2.013757
..	...	...
795	0.227143	1.474793
796	4.434784	-0.931543
797	1.449455	-0.170504
798	-2.888282	2.721490
799	-7.000656	2.805481

[800 rows x 2 columns]

```
[56]: plt.figure()
plt.scatter(df_PC['PC1'], df_PC['PC2'])
plt.xlabel('First PC')
plt.ylabel('Second PC')
plt.title('PCA')
plt.hlines(0, -10, 10, linestyle = 'dotted', colors = 'black')
plt.vlines(0, -10, 10, linestyle = 'dotted', colors = 'black')
```

```
[56]: <matplotlib.collections.LineCollection at 0x1f2bf554e88>
```



According to the tables above, the variables 'bi1', 'bi2', 'use3', 'use4', and 'pu3' appear most strongly correlated on the first principal component. 'peu1', 'incl', 'sa3', 'sal', and 'enj2' are the most strongly correlated on the second component.

7. (5 points) Calculate the proportion of variance explained (PVE) and the cumulative PVE for all the principal components. Approximately how much of the variance is explained by the first two principal components?

```
[57]: percentage1 = 100 * pca.explained_variance_ratio_[0]
percentage2 = 100 * pca.explained_variance_ratio_[1]
sum_per = 100 * (sum(pca.explained_variance_ratio_))
print(percentagel, 'percent of the variance is explained by the first principal_
↪component')
print(percentagel2, 'percent of the variance is explained by the second_
↪principal component')
print(sum_per, 'percent of the variance is explained by the two components_
↪together')
```

```
22.810627785663385 percent of the variance is explained by the first principal
component
6.37247429597453 percent of the variance is explained by the second principal
component
29.183102081637912 percent of the variance is explained by the two components
together
```

8. (10 points) Perform *t*-SNE on the dataset and plot the observations on the first and second dimensions. Describe your results.

```
[58]: tsne = TSNE(n_components = 2, random_state = 456123)
df_tsne = pd.DataFrame(data = tsne.fit_transform(df_wiki), columns = ['dim1', 'dim2'])
df_tsne
```

```
[58]:
```

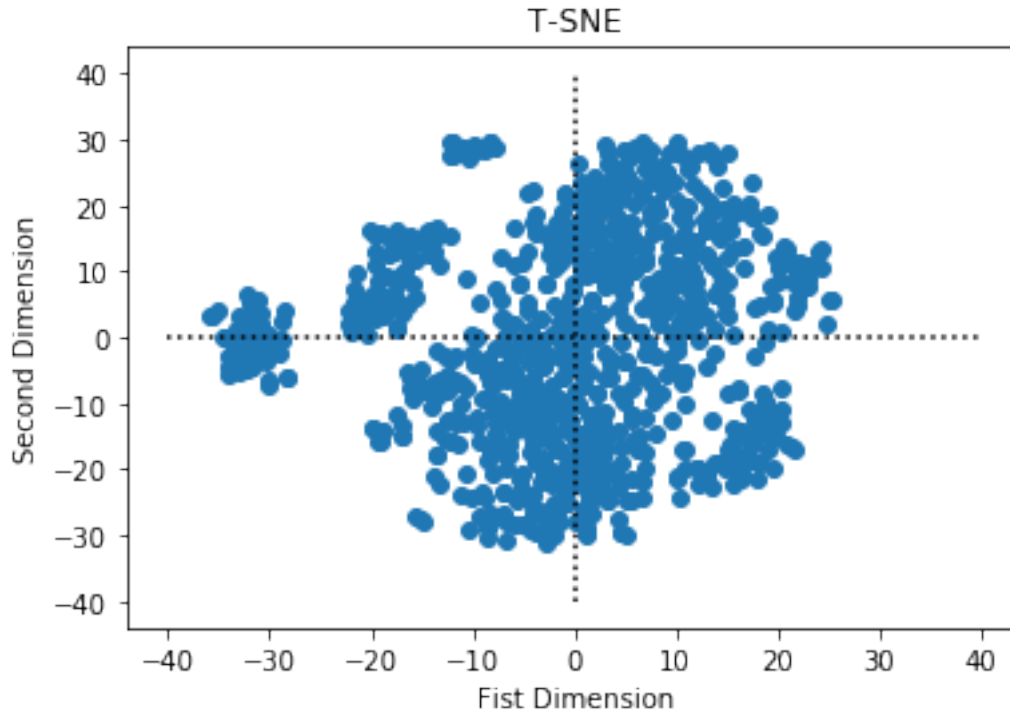
	dim1	dim2
0	-19.583126	11.007133
1	-15.840316	14.153754
2	-32.024673	5.357511
3	-33.440903	-4.613192
4	-10.368248	-28.905149
..	...	...
795	17.824806	-2.804761
796	6.833607	-19.981953
797	5.579870	-12.514900
798	1.697928	5.827499
799	12.196272	28.150850

[800 rows x 2 columns]

```
[59]: plt.figure()
plt.hlines(0, -40, 40, linestyle = 'dotted', colors = 'black')
plt.vlines(0, -40, 40, linestyle = 'dotted', colors = 'black')
plt.scatter(df_tsne['dim1'], df_tsne['dim2'])
plt.xlabel('First Dimension')
plt.ylabel('Second Dimension')
plt.title('T-SNE')
```

```
[59]: Text(0.5, 1.0, 'T-SNE')
```








According to the plot, we can see a large cluster exists in the middle and several small ones surround. The large cluster implies that there are a majority of features that are similar to one another.

9. (15 points) Perform  $k$ -means clustering with  $k = 2, 3, 4$ . Be sure to scale each feature (i.e., mean zero and standard deviation one). Plot the observations on the first and second principal components from PCA and color-code each observation based on their cluster membership. Discuss your results.

```
[60]: def plotting_kmean(k):
    k_mean = KMeans(n_clusters = k, random_state = 456123).fit(df_wiki)
    plt.figure()
    plt.scatter(df_PC['PC1'], df_PC['PC2'], c = k_mean.labels_)
    plt.xlabel('First PC')
    plt.ylabel('Second PC')
    plt.title(f'PCA based on {k}-means clustering')
    plt.hlines(0, -10, 10, linestyle = 'dotted', colors = 'black')
    plt.vlines(0, -10, 10, linestyle = 'dotted', colors = 'black')
```

```
[61]: k = [2,3,4]
    for i in k:
        plotting_kmean(i)
```






According to the graphs above, we can see that 2-means clustering performs the best as there are not many overlappings between the clusters. The 3-means clustering also perform well that we can observe clear boundaries for the clusters. For these two plots, the first PC has influential effects as the boundaries are quite vertical for 2-means and for 3-means. The boundaries for 3-means are little bit tilt, which suggest the second pc also has influence on clustering. The overlapping in the 4-means that the current two pcs are not enough to get four clusters.

10. (10 points) Use the elbow method, average silhouette, and/or gap statistic to identify the optimal number of clusters based on  $k$ -means clustering with scaled features.


```
[62]: #referenced from https://predictivehacks.com/
      ↪ k-means-elbow-method-code-for-python/
distortions = []
k = range(1,10)
for i in k:
    kmeanModel = KMeans(n_clusters = i)
    kmeanModel.fit(df_wiki)
    distortions.append(kmeanModel.inertia_)
plt.figure()
plt.plot(k, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The elbow method')
```

```
[62]: Text(0.5, 1.0, 'The elbow method')
```



```
[64]: #referenced from https://scikit-learn.org/stable/modules/generated/sklearn.  
      ↪metrics.silhouette_score.html  
sil = []  
for i in range(2,11):  
    kmeanModel = KMeans(n_clusters = i)  
    sil.append(silhouette_score(df_wiki,kmeanModel.fit_predict(df_wiki)))  
plt.figure()  
plt.plot(range(2,11), sil, 'bx-')  
plt.ylabel('Average Silhouette Score')  
plt.xlabel('Number of Clusters')  
plt.title('Average Silhouette')
```

```
[64]: Text(0.5, 1.0, 'Average Silhouette')
```




```
[73]: from gap_statistic import optimalK
      opt = optimalK.OptimalK()
      opt_clusters = opt(df_wiki, cluster_array = np.arange(2,11))
      print("the optimal number of cluster based on gap statistic is", opt_clusters)
```

the optimal number of cluster based on gap statistic is 10

According to the two graphs, the optimal number of clusters is 2.

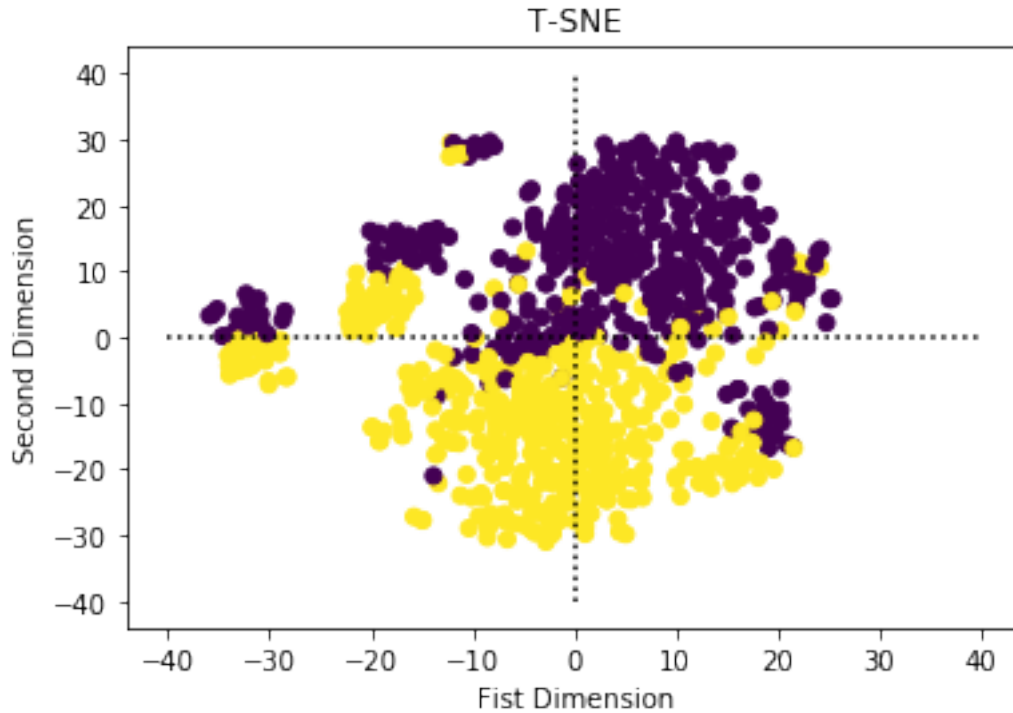
11. (15 points) Visualize the results of the optimal  $\hat{k}$ -means clustering model. First use the first and second principal components from PCA, and color-code each observation based on their cluster membership. Next use the first and second dimensions from  $t$ -SNE, and color-code each observation based on their cluster membership. Describe your results. How do your interpretations differ between PCA and  $t$ -SNE?

```
[70]: plotting_kmean(2)
```



```
[71]: tsne = TSNE(n_components = 2, random_state = 456123)
df_tsne = pd.DataFrame(data = tsne.fit_transform(df_wiki), columns = ['dim1',
↪ 'dim2'])
k2 = KMeans(n_clusters=2, random_state=456123).fit(df_wiki)
plt.figure()
plt.hlines(0, -40, 40, linestyle = 'dotted', colors = 'black')
plt.vlines(0, -40, 40, linestyle = 'dotted', colors = 'black')
plt.scatter(df_tsne['dim1'], df_tsne['dim2'], c = k2.labels_)
plt.xlabel('Fist Dimension')
plt.ylabel('Second Dimension')
plt.title('T-SNE')
```

```
[71]: Text(0.5, 1.0, 'T-SNE')
```



From the graphs above we can see that PCA does a better job in sperating the data in 2 cluster. By comparison, the t\_SNE has a lot overlaps for the two clusters. The reason why that t-SNE only preserves local similarities whereas PCA perserves large pairwise distance maximized variance. Since t-SNE only finds the local neighbors, it may not be able to get clear cluster boundaries. Moreover, the t-SNE finds patterns in the data by identify the similarity, while PCA focus on placing dissimilar data point far apart. Therefore, this can be another reason for why PCA does a better job at separating different clusters.